
Checkpoint 2

p7zip

Arnav Nidumolu, Atharva Kale, Pascal von Fintel, Patrick
Negus

2023-04-06

Checkpoint 2

Contents:

- [Checkpoint 2](#)
 - [Contents:](#)
 - [Static Analysis](#)
 - [Dynamic Analysis](#)
 - * [Generating a corpus](#)
 - * [Experimenting with fuzzing composition flags](#)
 - * [Parallel fuzzing](#)
 - * [Results](#)

Github Link

<https://github.com/atharvakale343/390r-debugging-setup>

Dynamic Analysis

Fuzzing

Fuzzing was the main dynamic analysis technique we used against our target `p7zip`. We mainly fuzzed the extract (`e`) feature of our binary as the feature uses several decompression algorithms as part of its execution.

We used `afl-plus-plus` as the primary fuzzing tool.

<https://github.com/AFLplusplus/AFLplusplus>

Generating a corpus

We took a variety of steps to find a good enough corpus for our fuzzing efforts. The major approach here was to search online for commonly used corpora. Our main goal here was to find not only `.zip` format, but also as many different formats possible.

We found a decent corpus at <https://github.com/strongcourage/fuzzing-corpus>

This included the following formats:

- `.zip`
- `.gzip`
- `.lrzip`
- `.jar`

We added this as a target to our fuzzing Makefile.

```
1 get-inputs:
2     rm -rf in_raw fuzzing-corpus && mkdir in_raw
3
4     git clone -n --depth=1 --filter=tree:0 git@github.com:strongcourage
      /fuzzing-corpus.git
5     cd fuzzing-corpus && git sparse-checkout set --no-cone zip gzip/go-
      fuzz lrzip jar && git checkout
6     mv fuzzing-corpus/zip/go-fuzz/* in_raw
7     mv fuzzing-corpus/jar/* in_raw
8     mv fuzzing-corpus/gzip/go-fuzz/* in_raw
9     mv fuzzing-corpus/lrzip/* in_raw
```

The next step was to choose only “interesting” inputs from this corpus. This includes small inputs that don’t crash that binary immediately.

We used the `afl-cmin` functionality to minimize the corpus.

```
1 afl-cmin -i in_raw -o in_unique -- $(BIN_AFL) e -y @@
```

Another important minimization step included `tmin`. This augments each input such that it can be as small as possible without compromising its ability to mutate and produce coverage in the instrumented target.

Unfortunately, this process takes a long time, and it only completed for us after a day.

```
1 cd in_unique; for i in *; do afl-tmin -i "$$i" -o "../in/$$i" -- ../$(  
    BIN_AFL) e -y @@; done
```

The cybersec room servers come in handy here!

Experimenting with fuzzing composition flags

We discovered that it is not enough to fuzz a plain instrumented target with `afl-plus-plus`. The target binary may not be easily crashed with mutated inputs as `p7zip` has a robust input error checker. We took to fuzzing with various sanitizers instead to search for harder to find bugs.

We used the following sanitizers on our target:

- ASAN: Address Sanitizer: discovers memory error vulnerabilities such as use-after-free, heap/buffer overflows, initialization order bugs etc.
- MSAN: Memory Sanitizer: mainly used to discover reads to uninitialized memory such as structs etc.
- TSAN: Thread Sanitizer: finds race conditions

```
1 fuzz-afl:  
2     AFL_SKIP_CPUFREQ=1 AFL_I_DONT_CARE_ABOUT_MISSING_CRASHES=1 $(  
        AFL_FUZZ) -M main-afl-$(HOSTNAME) -t 30000 -i in -o out -- $(  
        BIN_AFL) e -y @@  
3  
4 fuzz-afl-asan:  
5     AFL_SKIP_CPUFREQ=1 AFL_I_DONT_CARE_ABOUT_MISSING_CRASHES=1 $(  
        AFL_FUZZ) -S variant-afl-asan -t 30000 -i in -o out -- $(  
        BIN_AFL_ASAN) e -y @@  
6  
7 fuzz-afl-msan:  
8     AFL_SKIP_CPUFREQ=1 AFL_I_DONT_CARE_ABOUT_MISSING_CRASHES=1 $(  
        AFL_FUZZ) -S variant-afl-msan -t 30000 -i in -o out -- $(  
        BIN_AFL_MSAN) e -y @@
```

```
9
10 fuzz-afl-tsan:
11     AFL_SKIP_CPUFREQ=1 AFL_I_DONT_CARE_ABOUT_MISSING_CRASHES=1 $(
        AFL_FUZZ) -S variant-afl-tsan -t 30000 -i in -o out -- $(
        BIN_AFL_TSAN) e -y @@
```

Parallel Fuzzing

To start with, our approach was to fuzz the `extract` command of `7zz`. So we found an appropriate corpus and fuzzed with the `e` command-line argument (along with `-y` to account for same filenames / avoid user input hangs).

With all different sets of compilation flags that we mentioned previously, we compiled the binaries with AFL instrumentation. Then, to more effectively fuzz, we setup a parallel fuzzing environment in one of the **CyberSec club** VMs.

We added the `afl-fuzz` commands in a `Makefile` and followed the official [guide](#) for using multiple cores. Below are the commands we utilized. All of our fuzzers shared the same input and output directories to keep track of current fuzzing state.

```
1 AFL_SKIP_CPUFREQ=1 AFL_I_DONT_CARE_ABOUT_MISSING_CRASHES=1 $(AFL_FUZZ)
  -M main-afl-$
2 (HOSTNAME) -t 2000 -i in -o out -- $(BIN_AFL) e -y @@
```

Our main fuzzer used a regular instrumented AFL binary with no other `CFLAGS`. We used a timeout of 2 seconds to denote a hang (or infinite loops).

```
1 AFL_SKIP_CPUFREQ=1 AFL_I_DONT_CARE_ABOUT_MISSING_CRASHES=1 $(AFL_FUZZ)
  -S variant-af
2 l-asan -t 2000 -i in -o out -- $(BIN_AFL_ASAN) e -y @@
```

Our variant fuzzers utilized binaries compiled with other flags (such as `asan` and `msan`). These had the same timeout as before of 2 seconds.

To keep track of all fuzzers and run them simultaneously, we used `tmux` sessions with a separate window for each fuzzer.

Results

We ran the fuzzers using multiple cores for around 2.5 days. We noticed no crashes in most of the variants, with `msan` being the exception. However, some fuzzers encountered hangs.

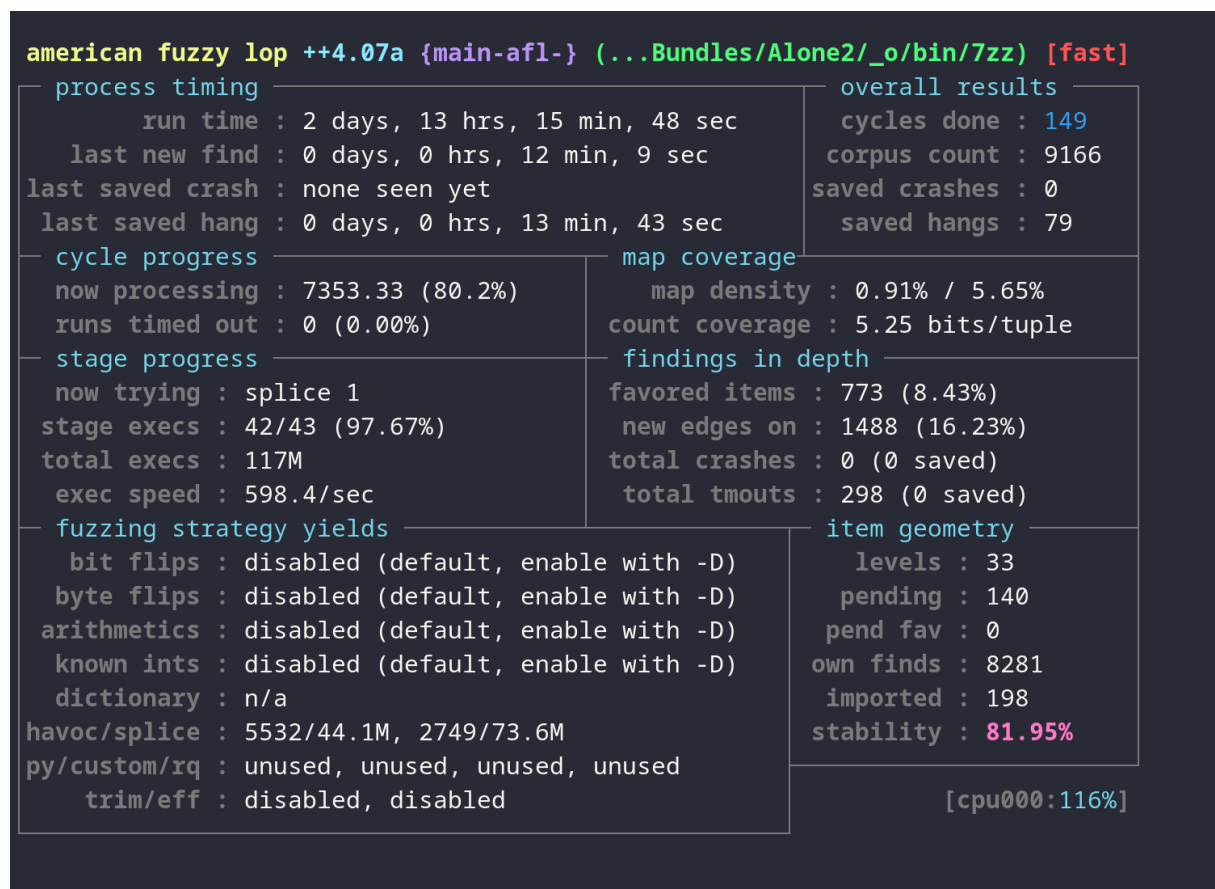



Figure 1: Main AFL Fuzzer





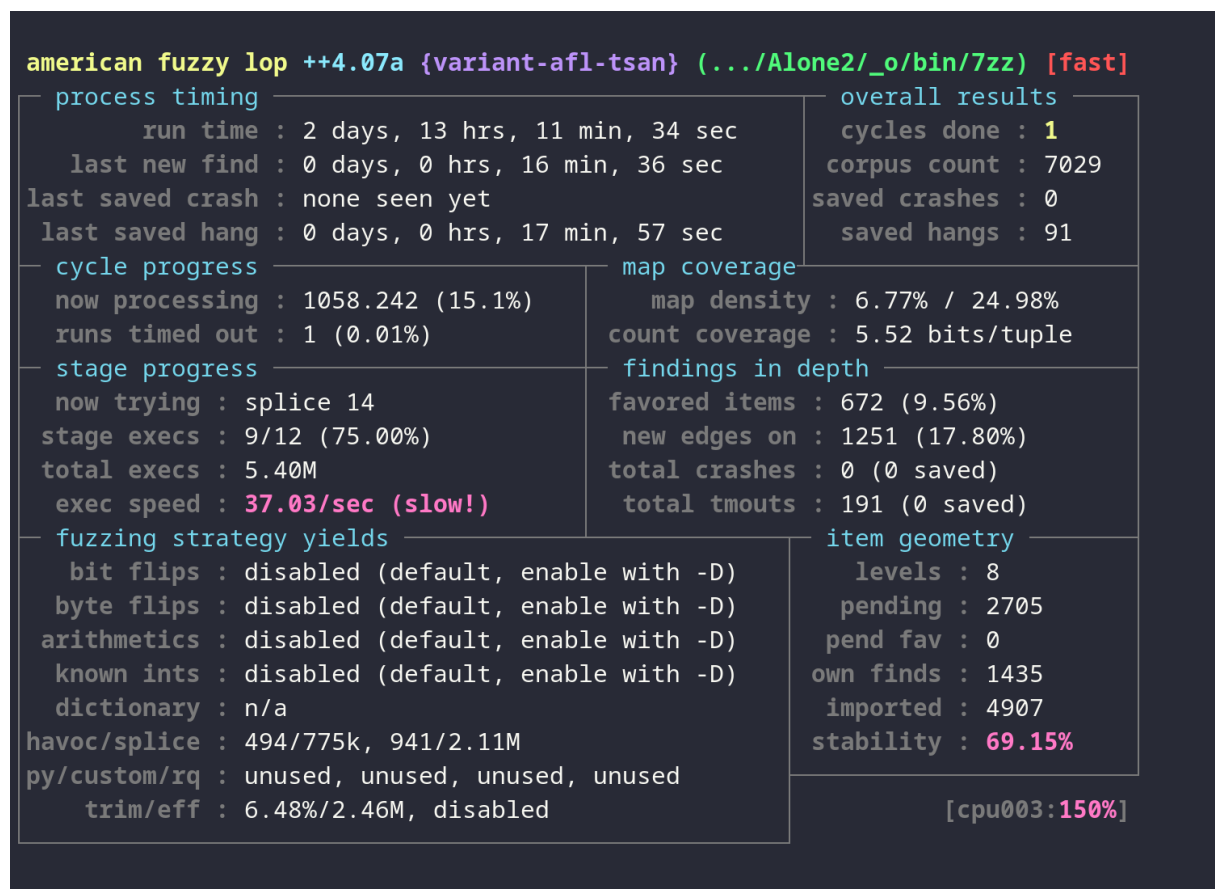


Figure 4: Tsan Variant Fuzzer

We tried running an input from `in/hangs` to check where an infinite loop could occur. But, all inputs eventually completed while taking longer than 2 seconds, so we concluded that the timeout value was too low. These executions were incorrectly flagged as hangs due to relatively low timeouts. For our next fuzzing attempts, we plan to increase this and make the timeout around 30 seconds to account for larger file inputs.

Analyzing msan crashes

As the msan variant was the only one that produced crashes, we compiled a msan binary with debug flags and analyzed the crash.

```

+ fuzzing-work git:(main) x ../../7zz_afl_msan_dbg/CPP/7zip/Bundles/Alone2/_o/bin/7zz e -y in/signed.jar [15/1981]
7-Zip (z) 22.00 ZS v1.5.2 (x64) : Copyright (c) 1999-2022 Igor Pavlov : 2022-06-15
64-bit locale=en_US.UTF-8 Threads:6

Scanning the drive for archives:
1 file, 1781 bytes (2 KiB)

Extracting archive: in/signed.jar
--
Path = in/signed.jar
Type = zip
Physical Size = 1781

0%==3949344==WARNING: MemorySanitizer: use-of-uninitialized-value
#0 0x55555575245e in NWindows::NFile::NDir::SetFileAttrib_PosixHighDetect(char const*, unsigned int) /home/ubuntu/390r-debugging-setup/7zz_afl_msan_dbg/CPP/7zip/Bundles/Alone2/../../../../Windows/FileDir.cpp:1085:12
#1 0x555555e23629 in CArchiveExtractCallback::SetAttrib() /home/ubuntu/390r-debugging-setup/7zz_afl_msan_dbg/CPP/7zip/Bundles/Alone2

```

Figure 5: crash output from msan

As displayed above, the crash occurs due to a use of uninitialized value in *FileDir.cpp*.

```

1085 else if (S_ISLNK(st.st_mode))
1086 {
1087     /* for most systems: permissions for symlinks are fixed to rwxrwxrwx.
1088     so we don't need chmod() for symlinks. */
1089     return true;
1090     // SetLastError(ENOSYS);
}

[ STACK ]
00:0000 rsp 0x7fffffffef60 ← 0x0
01:0008 rdx 0x7fffffffef68 ← 0x803
02:0010 0x7fffffffef70 ← 0x275c09 /* "\t\t\t" */
03:0018 0x7fffffffef78 ← 0x1
04:0020 0x7fffffffef80 → 0x3e8000081b4 ← 0x0
05:0028 0x7fffffffef88 ← 0x3e8
06:0030 0x7fffffffef90 ← 0x0
07:0038 0x7fffffffef98 ← 0x79 /* 'y' */

[ BACKTRACE ]
f 0 0x55555575229c NWindows::NFile::NDir::SetFileAttrib_PosixHighDetect(char const*, unsigned int)+268
f 1 0x555555e2362a CArchiveExtractCallback::SetAttrib()+970
f 2 0x555555e2a609 CArchiveExtractCallback::SetOperationResult(int)+1817
f 3 0x5555555bc962f NArchive::NZip::CHandler::Extract(unsigned int const*, unsigned int, int, IArchiveExtractCallback*)+7775
f 4 0x555555e80530 Extract(CCodecs*, CObjectVector<COpenType> const&, CRecordVector<int> const&, CObjectVector<UString>&, CObjectVector<UString> const&, CExtractorOptions const&, IOpenCallbackUI*, IExtractCallbackUI*, IHashCalc*, UString&, CDecompressStat&)+20704
f 5 0x555555e80530 Extract(CCodecs*, CObjectVector<COpenType> const&, CRecordVector<int> const&, CObjectVector<UString>&, CObjectVector<UString> const&, CExtractorOptions const&, IOpenCallbackUI*, IExtractCallbackUI*, IHashCalc*, UString&, CDecompressStat&)+20704
f 6 0x555555f3bc69 Main2(int, char**)+26233
f 7 0x555555f45b75 main+213

pwndbg> p st.st_mode
$1 = 33204
pwndbg>

```

Figure 6: gdb analysis for msan crash

From previous error message, we see that msan has flagged `st.st_mode` as uninitialized. But, looking into **gdb**, this variable seems to be defined, set to 33204. This is because it was called with `lstat(path, &st)` at the beginning of the function, which initialized all fields of the struct.

From this, we can conclude that msan had incorrectly flagged this an uninitialized and this is a *false positive*. All other msan crashes refer to the same line, so we determined that **msan** is not a good fit for this project, possibly missing out on initializers.

We looked more into [clang documentation](#), which affirmed this belief:

it may introduce false positives and therefore should be used with care

Static Analysis

Codeql

To analyze the code for common C/C++ bugs, we used **codeql** to scan the source code.

We first created a analysis database by providing the **make** instructions to codeql.

```
1 cd p7zip/CPP/7zip/Bundles/Alone2
2 codeql database create ../../../../codeql-playground/analysis-db.
  codeql -l cpp -c "make -B -f makefile.gcc" --overwrite
3 cd -
```

Then, we download *cpp-queries* and tested the produced database against it.

```
1 codeql pack download codeql/cpp-queries
2 codeql database analyze analysis-db.codeql --format CSV --output
  analysis.csv
```

```
Starting evaluation of codeql/cpp-queries/Security/CWE/CWE-704/WcharCharConversion.q1.
[42/47 eval 8ms] Evaluation done; writing results to codeql/cpp-queries/Security/CWE/CWE-676/DangerousUseOfCin.bqrs.
Starting evaluation of codeql/cpp-queries/Security/CWE/CWE-732/OpenCallMissingModeArgument.q1.
[43/47 eval 63ms] Evaluation done; writing results to codeql/cpp-queries/Security/CWE/CWE-704/WcharCharConversion.bqrs.
Starting evaluation of codeql/cpp-queries/Security/CWE/CWE-732/UnsafeDaclSecurityDescriptor.q1.
[44/47 eval 28ms] Evaluation done; writing results to codeql/cpp-queries/Security/CWE/CWE-732/OpenCallMissingModeArgument.bqrs.
Starting evaluation of codeql/cpp-queries/Summary/LinesOfCode.q1.
[45/47 eval 13ms] Evaluation done; writing results to codeql/cpp-queries/Security/CWE/CWE-732/UnsafeDaclSecurityDescriptor.bqrs.
Starting evaluation of codeql/cpp-queries/Summary/LinesOfUserCode.q1.
[46/47 eval 4ms] Evaluation done; writing results to codeql/cpp-queries/Summary/LinesOfCode.bqrs.
[47/47 eval 2.1s] Evaluation done; writing results to codeql/cpp-queries/Summary/LinesOfUserCode.bqrs.
Shutting down query evaluator.
Interpreting results.
Analysis produced the following diagnostic data:
| Diagnostic | Summary |
|-----+-----+
| Successfully extracted files | 44 results |

Analysis produced the following metric data:
| Metric | Value |
|-----+-----+
| Total lines of user written C/C++ code in the database | 4129 |
| Total lines of C/C++ code in the database | 400756 |

→ codeql-playground git:(main) ✗ cat analysis.csv
→ codeql-playground git:(main) ✗
```

Figure 7: Codesql statistics

But, this did not output any glaring errors, executing without any warnings or useful metrics. We concluded that codeql only utilizes simple checks which would have already been accounted for in the

source code.

CPPCheck

We ran the codebase through the static analysis tool cppcheck, which tagged 1569 warnings and errors. One of the common errors flagged by cppcheck was shiftTooManyBits

[390r-debugging-setup\p7zip\CPP\7zip\Archive\7z\7zIn.cpp](#)

261	shiftTooManyBits	758	error	Shifting 32-bit value by 32 bits is undefined behaviour
1546	shiftTooManyBits	758	error	Shifting 32-bit value by 32 bits is undefined behaviour
1547	shiftTooManyBits	758	error	Shifting 32-bit value by 32 bits is undefined behaviour
1598	shiftTooManyBits	758	error	Shifting 32-bit value by 62 bits is undefined behaviour

Unfortunately, when looking at the actual source code, almost all of these errors come from an innocuous function:

```
#define GetUi64(p) (GetUi32(p) | ((UInt64)GetUi32(((const Byte *)(p)) + 4) << 32))

#define GetUi32(p) ( \
    ((const Byte *)(p))[0] | \
    ((UInt32)((const Byte *)(p))[1] << 8) | \
    ((UInt32)((const Byte *)(p))[2] << 16) | \
    ((UInt32)((const Byte *)(p))[3] << 24))
```

The rest, on closer inspection, are also falsely flagged as errors, such as this one:

2594	shiftTooManyBits	758	error	Shifting 32-bit value by 63 bits is undefined behaviour
----------------------	------------------	---------------------	-------	---

```
2594 if (node.FileSize ≥ ((UInt64)1 << 63))
2595 return S_FALSE;
```

A more promising error seems to be a possible null pointer exception:

[390r-debugging-setup\p7zip\CPP\7zip\Archive\Zip\ZipHandlerOut.cpp](#)

41	nullPointerRedundantCheck	476	warning	Either the condition 'password' is redundant or there is possible null pointer dereference: s++.
41	nullPointerArithmeticRedundantCheck	682	warning	Either the condition 'password' is redundant or there is pointer arithmetic with NULL pointer.

```
37 static bool IsSimpleAsciiString(const wchar_t *s)
38 {
39     for (;;)
40     {
41         wchar_t c = *s++;
42         if (c == 0)
43             return true;
44         if (c < 0x20 || c > 0x7F)
45             return false;
46     }
47 }
```

This function is only called once, in the same file at line 415:

```
392 CMyComPtr<ICryptoGetTextPassword2> getTextPassword;
393 {
394     CMyComPtr<IArchiveUpdateCallback> udateCallBack2(callback);
395     udateCallBack2.QueryInterface(IID_ICryptoGetTextPassword2, &getTextPassword);
396 }
397 CCompressionMethodMode options;
398 (CBaseProps &)options = _props;
399 options._dataSizeReduce = largestSize;
400 options._dataSizeReduceDefined = largestSizeDefined;
401
402 options.PasswordIsDefined = false;
403 options.Password.Wipe_and_Empty();
404 if (getTextPassword)
405 {
406     CMyComBSTR_Wipe password;
407     Int32 passwordIsDefined;
408     RINOK(getTextPassword->CryptoGetTextPassword2(&passwordIsDefined, &password));
409     options.PasswordIsDefined = IntToBool(passwordIsDefined);
410     if (options.PasswordIsDefined)
411     {
412         if (!m_ForceAesMode)
413             options.IsAesMode = thereAreAesUpdates;
414
415         if (!IsSimpleAsciiString(password))
416             return E_INVALIDARG;
```

It looks like `password` gets populated in `CryptoGetTexPassword2`, looking at that function, and the subsequent call to `StringToBstr`, it unfortunately looks like the nullpointer is properly checked for.

```
97  STDMETHODIMP CUpdateCallback100Imp::CryptoGetTextPassword2(Int32 *passwordIsDefined, BSTR *password)
98  {
99      *password = NULL;
100      *passwordIsDefined = BoolToInt(PasswordIsDefined);
101      if (!PasswordIsDefined)
102          return S_OK;
103      return StringToBstr(Password, password);
104  }

77  inline HRESULT StringToBstr(LPCOLESTR src, BSTR *bstr)
78  {
79      *bstr = ::SysAllocString(src);
80      return (*bstr) ? S_OK : E_OUTOFMEMORY;
81  }
```