
Final Writeup

p7zip

Arnav Nidumolu, Atharva Kale, Pascal von Fintel, Patrick
Negus

2023-05-14

Final Writeup

Public Github Repository - This should include all code you wrote for eg. static analysis, fuzzing harnesses, etc. If you built your target with instrumentation for the purposes of fuzzing, this should also include build scripts. If you performed reverse engineering on your target and eg. started renaming variables/functions/did work on that front, include the relevant ghidra files as well.

Start your writeup with a description of what you learned about this target. This should include some notes about the code layout, maybe some coding practices you noticed while going through the target or just more general functionality. Which parts of the target did you think were most interesting for the purposes of finding bugs?

Describe what you chose for your automated analysis portion and why. How did you set this up, did you encounter issues (eg. slow fuzzer performance), and if so what did you do to improve on these issues.

What were the biggest challenges you faced when dealing with your target?

If given more time, what do you think would be good next steps to continue doing research on the target with the goal of finding bugs?

Contents:

- [Github Repository](#)
- [Overview of the Target](#)
 - [Code Layout](#)
 - [Coding Observations](#)
 - [Analyzing Target Features](#)
- [Automated Analysis](#)
 - [Fuzzing](#)
 - * How was it set up
 - * Results etc...
 - [Static Analysis](#)
 - * ...
- [Challenges Faced](#)
 - ...
- [Next Steps](#)

Github Link

<https://github.com/atharvakale343/p7zip-390r>

Overview of the Target

p7zip is a fully compliant linux port of the open source *7zip* tool for Windows. It is a utility used to archive and extract various compression formats. It is primarily used in Windows GUI tools as an underlying utility to support their file compression features.

p7zip provides the following features:

1. Several compression algorithms (*lz4*, *zstd*, *Lizard*, etc...)
2. CLI frontend
3. Cryptographic algorithms for Home Exp archive encryption (*SHA256*, *AES*, *RAR5*, etc...)

Code Layout

[Call Graph for extract command](#)

[Call Graph for archive command](#)

Coding Observations

Target Features

The main features of 7zz is to extract from different compression formats and archive a collection of files. We decided to focus mainly on the [extract](#) or **e** command which accepts a compressed file and extracts its contents to the current directory.

We also looked into the [archive](#) or **a** command that compresses a list of files into a .7z file. This uses various compression algorithms such as LZ4, Brotli, Lizard, etc.

Automated Analysis

Fuzzing

Fuzzing was the main dynamic analysis technique we used against our target [p7zip](#). We mainly fuzzed the extract (**e**) feature of our binary as the feature uses several decompression algorithms as part of its execution.

We used [afl-plus-plus](#) as the primary fuzzing tool.

<https://github.com/AFLplusplus/AFLplusplus>

Generating a corpus

We took a variety of steps to find a good enough corpus for our fuzzing efforts. The major approach here to was to search online for commonly used corpora. We wanted to find not only .zip format, but also as many different formats possible.

We found a decent corpus at <https://github.com/strongcourage/fuzzing-corpus>

This included the following formats:

- .zip
- .gzip
- .lrzip

- `.jar`

We added this as a target to our fuzzing Makefile.

```
1 get-inputs:
2     rm -rf in_raw fuzzing-corpus && mkdir in_raw
3
4     git clone -n --depth=1 --filter=tree:0 git@github.com:strongcourage
      /fuzzing-corpus.git
5     cd fuzzing-corpus && git sparse-checkout set --no-cone zip gzip/go-
      fuzz lrzip jar && git checkout
6     mv fuzzing-corpus/zip/go-fuzz/* in_raw
7     mv fuzzing-corpus/jar/* in_raw
8     mv fuzzing-corpus/gzip/go-fuzz/* in_raw
9     mv fuzzing-corpus/lrzip/* in_raw
```

The next step was to choose only “interesting” inputs from this corpus. This includes small inputs that don’t crash that binary immediately.

We used the `afl-cmin` functionality to minimize the corpus.

```
1 afl-cmin -i in_raw -o in_unique -- $(BIN_AFL) e -y @@
```

Another important minimization step included `tmin`. This augments each input such that it can be as small as possible without compromising its ability to mutate and produce coverage in the instrumented target.

Unfortunately, this process takes a long time, and it only completed for us after a day.

```
1 cd in_unique; for i in *; do afl-tmin -i "$$i" -o "../in/$$i" -- ../$(
      BIN_AFL) e -y @@; done
```

The cybersec room servers come in handy here!

Experimenting with fuzzing composition flags

We discovered that it is not enough to fuzz a plain instrumented target with `afl-plus-plus`. The target binary may not be easily crashed with mutated inputs as `p7zip` has a robust input error checker. We took to fuzzing with various sanitizers instead to search for harder to find bugs.

We used the following sanitizers on our target:

- ASAN: Address Sanitizer: discovers memory error vulnerabilities such as use-after-free, heap/buffer overflows, initialization order bugs etc.
- MSAN: Memory Sanitizer: mainly used to discover reads to uninitialized memory such as structs etc.

- TSAN: Thread Sanitizer: finds race conditions

```
1 afl:
2   rm -rf $(BIN_AFL)
3   git clone $(GH_URL) $(BIN_AFL)
4   cp 7zz-makefiles/$(BIN_DEFAULT).mak $(BIN_AFL)/CPP/7zip/7zip_gcc.
   mak
5   cd $(BIN_AFL)/CPP/7zip/Bundles/Alone2 && CC=$(AFL_CC) CXX=$(AFL_CXX)
   make -f makefile.gcc
6
7 afl-asan:
8   rm -rf $(BIN_AFL_ASAN)
9   git clone $(GH_URL) $(BIN_AFL_ASAN)
10  cp 7zz-makefiles/$(BIN_AFL_ASAN).mak $(BIN_AFL_ASAN)/CPP/7zip/7
   zip_gcc.mak
11  cd $(BIN_AFL_ASAN)/CPP/7zip/Bundles/Alone2 && AFL_USE_ASAN=1 CC=$(
   AFL_CC) CXX=$(AFL_CXX) make -f makefile.gcc
12
13 afl-msan:
14   rm -rf $(BIN_AFL_MSAN)
15   git clone $(GH_URL) $(BIN_AFL_MSAN)
16   cp 7zz-makefiles/$(BIN_AFL_MSAN).mak $(BIN_AFL_MSAN)/CPP/7zip/7
   zip_gcc.mak
17   cd $(BIN_AFL_MSAN)/CPP/7zip/Bundles/Alone2 && AFL_CC_COMPILER=LLVM
   AFL_USE_MSAN=1 CC=$(AFL_CC) CXX=$(AFL_CXX) make -f makefile.gcc
18
19 afl-tsan:
20   rm -rf $(BIN_AFL_TSAN)
21   git clone $(GH_URL) $(BIN_AFL_TSAN)
22   cp 7zz-makefiles/$(BIN_AFL_TSAN).mak $(BIN_AFL_TSAN)/CPP/7zip/7
   zip_gcc.mak
23   cd $(BIN_AFL_TSAN)/CPP/7zip/Bundles/Alone2 && AFL_USE_TSAN=1 CC=$(
   AFL_CC) CXX=$(AFL_CXX) make -f makefile.gcc
```

Extract command

Parallel Fuzzing

To start with, our approach was to fuzz the `extract` command of `7zz`. So we found an appropriate corpus and fuzzed with the `e` command-line argument (along with `-y` to account for same filenames / avoid user input hangs).

With all different sets of compilation flags that we mentioned previously, we compiled the binaries with AFL instrumentation. Then, to more effectively fuzz, we setup a parallel fuzzing environment in one of the **CyberSec club** VMs.

We added the `afl-fuzz` commands in a `Makefile` and followed the official [guide](#) for using multiple

cores. Below are the commands we utilized. All of our fuzzers shared the same input and output directories to keep track of current fuzzing state.

```
1 AFL_SKIP_CPUFREQ=1 AFL_I_DONT_CARE_ABOUT_MISSING_CRASHES=1 $(AFL_FUZZ)
  -M main-afl-$(HOSTNAME) -t 2000 -i in -o out -- $(BIN_AFL) e -y @@
```

Our main fuzzer used a regular instrumented AFL binary with no other CFLAGS. We used a timeout of 30 seconds to denote a hang (or infinite loops).

```
1 AFL_SKIP_CPUFREQ=1 AFL_I_DONT_CARE_ABOUT_MISSING_CRASHES=1 $(AFL_FUZZ)
  -S variant-afl-asan -t 2000 -i in -o out -- $(BIN_AFL_ASAN) e -y @@
```

Our variant fuzzers utilized binaries compiled with other flags (such as *asan* and *msan*). These had the same timeout as before of 30 seconds.

To keep track of all fuzzers and run them simultaneously, we used *tmux* sessions with a separate window for each fuzzer.

Results

We ran the fuzzers using multiple cores for around 5 days. We noticed no crashes in most of the variants, with ASAN being the exception. However, some fuzzers encountered hangs.

```
american fuzzy lop ++4.07a {main-afl-} (...Bundles/Alone2/_o/bin/7zz) [fast]
├── process timing
│   ├── run time : 5 days, 0 hrs, 16 min, 37 sec
│   ├── last new find : 0 days, 0 hrs, 0 min, 3 sec
│   ├── last saved crash : none seen yet
│   └── last saved hang : 0 days, 0 hrs, 13 min, 57 sec
├── cycle progress
│   ├── now processing : 11.3k.0 (99.7%)
│   ├── runs timed out : 0 (0.00%)
│   └── stage progress
│       ├── now trying : havoc
│       ├── stage execs : 4446/8000 (55.58%)
│       ├── total execs : 207M
│       └── exec speed : 667.6/sec
├── fuzzing strategy yields
│   ├── bit flips : disabled (default, enable with -D)
│   ├── byte flips : disabled (default, enable with -D)
│   ├── arithmetics : disabled (default, enable with -D)
│   ├── known ints : disabled (default, enable with -D)
│   ├── dictionary : n/a
│   ├── havoc/splice : 6716/76.6M, 4317/130M
│   ├── py/custom/rq : unused, unused, unused, unused
│   └── trim/eff : disabled, disabled
├── map coverage
│   ├── map density : 1.01% / 6.76%
│   ├── count coverage : 5.10 bits/tuple
│   └── findings in depth
│       ├── favored items : 1016 (8.97%)
│       ├── new edges on : 1830 (16.16%)
│       ├── total crashes : 0 (0 saved)
│       └── total tmouts : 293 (0 saved)
├── item geometry
│   ├── levels : 44
│   ├── pending : 558
│   ├── pend fav : 9
│   ├── own finds : 11.0k
│   ├── imported : 117
│   └── stability : 87.39%
└── overall results
    ├── cycles done : 171
    ├── corpus count : 11.3k
    ├── saved crashes : 0
    └── saved hangs : 40
[cpu000: 83%]
```

Figure 1: Main AFL Fuzzer

```
american fuzzy lop ++4.07a {variant-afl-asan} (.../Alone2/_o/bin/7zz) [fast]
├── process timing
│   ├── run time : 5 days, 0 hrs, 20 min, 29 sec
│   ├── last new find : 0 days, 0 hrs, 5 min, 0 sec
│   ├── last saved crash : 0 days, 3 hrs, 27 min, 15 sec
│   └── last saved hang : 0 days, 0 hrs, 14 min, 18 sec
├── cycle progress
│   ├── now processing : 7313.63 (89.7%)
│   ├── runs timed out : 0 (0.00%)
│   └── stage progress
│       ├── now trying : splice 15
│       ├── stage execs : 5/12 (41.67%)
│       ├── total execs : 9.02M
│       └── exec speed : 35.04/sec (slow!)
├── fuzzing strategy yields
│   ├── bit flips : disabled (default, enable with -D)
│   ├── byte flips : disabled (default, enable with -D)
│   ├── arithmetics : disabled (default, enable with -D)
│   ├── known ints : disabled (default, enable with -D)
│   ├── dictionary : n/a
│   ├── havoc/splice : 616/1.56M, 1282/4.63M
│   ├── py/custom/rq : unused, unused, unused, unused
│   └── trim/eff : 6.78%/2.75M, disabled
├── map coverage
│   ├── map density : 5.57% / 27.89%
│   ├── count coverage : 5.67 bits/tuple
│   └── findings in depth
│       ├── favored items : 843 (10.34%)
│       ├── new edges on : 1578 (19.36%)
│       ├── total crashes : 11.4k (289 saved)
│       └── total tmouts : 25 (0 saved)
├── item geometry
│   ├── levels : 9
│   ├── pending : 3640
│   ├── pend fav : 3
│   ├── own finds : 1617
│   ├── imported : 6369
│   └── stability : 61.35%
└── overall results
    ├── cycles done : 3
    ├── corpus count : 8151
    ├── saved crashes : 289
    └── saved hangs : 11
[cpu001: 50%]
```

Figure 2: ASAN Variant Fuzzer

We tried running an input from `in/hangs` to check where an infinite loop could occur. But, all inputs

eventually terminated while taking longer than 30 seconds. Therefore, we concluded that these executions were incorrectly flagged as hangs due to large size of the file. We could possibly set the timeout even higher to avoid this issue.

Analyzing asan crashes

```
Extracting archive: vuln.zip
=====
==39559==ERROR: AddressSanitizer: requested allocation size 0x154ac771c7ffffff (0x154ac771c8001000 after adjustments for alignment, red
zones etc.) exceeds maximum supported size of 0x10000000000 (thread T0)
    #0 0x7fa6c50d9f98 in operator new[](unsigned long) (/lib64/libasan.so.8+0xd9f98) (BuildId: d34ee7d544aadb28cdcd8dd99198ee7130a1fe4d)
    #1 0xf4937d (/home/lifewhiz/projects/revEng/390r-debugging-setup/7zz_af1_asan/CPP/7zip/Bundles/Alone2/_o/bin/7zz+0xf4937d) (BuildId
: 30fd98054d6b9d16484801a99d3e4e5e9e187262)
    #2 0xf49868 (/home/lifewhiz/projects/revEng/390r-debugging-setup/7zz_af1_asan/CPP/7zip/Bundles/Alone2/_o/bin/7zz+0xf49868) (BuildId
: 30fd98054d6b9d16484801a99d3e4e5e9e187262)
    #3 0xedfe25 (/home/lifewhiz/projects/revEng/390r-debugging-setup/7zz_af1_asan/CPP/7zip/Bundles/Alone2/_o/bin/7zz+0xedfe25) (BuildId
: 30fd98054d6b9d16484801a99d3e4e5e9e187262)
==39559==HINT: if you don't care about these errors you may set allocator_may_return_null=1
SUMMARY: AddressSanitizer: allocation-size-too-big (/lib64/libasan.so.8+0xd9f98) (BuildId: d34ee7d544aadb28cdcd8dd99198ee7130a1fe4d) in
operator new[](unsigned long)
==39559==ABORTING
→ fuzzing-work git:(main) █
```

Figure 3: ASAN error output

Upon executing the ASAN compiled binary on one of the crash inputs, we found that it occurred due to “request allocation size exceeding maximum supported size”. This likely occurred due to malloc being called with a huge size argument and returning `NULL`.

```
HRESULT CUnpacker::UnpackData(IInStream *inStream,
    const CResource &resource, const CHeader &header,
    const CDatabase *db,
    CByteBuffer &buf, Byte *digest)
{
    // if (resource.IsSolid()) return E_NOTIMPL;

    UInt64 unpackSize64 = resource.UnpackSize;
    if (db)
        unpackSize64 = db->Get_UnpackSize_of_Resource(resource);

    size_t size = (size_t)unpackSize64;
    if (size != unpackSize64)
        return E_OUTOFMEMORY;

    buf.Alloc(size);

    CBufPtrSeqOutStream *outStreamSpec = new CBufPtrSeqOutStream();
    CMyComPtr<ISequentialOutStream> outStream = outStreamSpec;
    outStreamSpec->Init((Byte *)buf, size);

    return Unpack(inStream, resource, header, db, outStream, NULL, digest);
}
```

Figure 4: Location of error

The crash occurs at the `buf.Alloc` call, which executes a C++ **new** operation that internally calls `malloc`. Here, an argument of `unpackSize64` is passed into the function.

```

46
47 void Alloc(size_t size)
48 {
49     if (size != _size)
50     {
51         Free();
52         if (size != 0)
53         {
54             _items = new T[size];
55         }
56     }
57 }
58
59 [ STACK ]
00:0000 rsp 0x7fffffffbb70 ← 0x154ac771c7ffffff
01:0008 0x7fffffffbb78 → 0x7fffffffd00 ← 0x0
02:0010 rbp 0x7fffffffbb80 → 0x7fffffffbfb0 → 0x7fffffffbcb0 → 0x7fffffffbcd0 → 0x7ffffffbf50 ← ...
03:0018 0x7fffffffbb88 → 0x578e16 ← mov edi, 0x28
04:0020 0x7fffffffbb90 → 0x7fffffffd00 ← 0x0
05:0028 0x7fffffffbb98 → 0x76d2a8 ← 0x0
06:0030 0x7fffffffbbba0 → 0x7fffffffd80 ← 0x4034b5000010ae7
07:0038 0x7fffffffbbba8 → 0x7fffffffd00 ← 0x800000000
60
61 [ BACKTRACE ]
62
63 f 0 0x45bc1c
64 f 1 0x578e16
65 f 2 0x57abcd
66 f 3 0x56ba3f
67 f 4 0x64612e
68 f 5 0x6495ab
69 f 6 0x64a3ab
70 f 7 0x64a9f1 CArc::OpenStreamOrFile(COpenOptions&)+363
71
72 pwndbg> p/x size
73 $1 = 0x154ac771c7ffffff
74 pwndbg>

```

Figure 5: Analyzing size argument in GDB

As shown above, the `unpackSize64` argument is a large unsigned integer, so `malloc` fails to allocate this memory and ASAN instigates a crash. If we can control this size argument, this could be a potential bug.

```

→ fuzzing-work git:(main) xxd -p vuln.zip | head | grep ffffffff771c74a15
0300ffffffffffc771c74a1514320000003f0030000004003000400030000000
→ fuzzing-work git:(main)

```

Figure 6: Size in File header

We analyzed the input file in more detail and found `unpackSize64` (in little endian) within the header of the file. So we can attempt to modify this offset within the header and control the amount of memory malloc'd. But, this is not an outright segfault since C++ error handling accounts for this and throws an exception, which is caught by the `p7zip` error handler.

However, this is a potential bug if combined with static/taint analysis so see if we could perform a possible overflow due to some arithmetic operations performed on this size argument.

Archive command

We also fuzzed the `archive` command of `7zz`. For this, we initially chose a corpus of `.txt` files, fuzzing the `a` command-line argument (along with `-y` to avoid user input hangs). But, we could not directly fuzz this since `afl-fuzz` only supports one cli argument and `a` can be used with multiple files with the following syntax:

```
1 7zz a files.zip file1.txt file2.txt file3.txt
```

We concluded that it would not be sufficient to just archive one file so we decided to create a harness which would allow for multiple file as arguments.

Harness

Our approach for the harness is as follows:

- As `afl-fuzz` allows for only one input to the target binary, our harness would accept one file name as argument.
- Contents of this input file would be divided into chunks of 1000 bytes and one new file will be created for each chunk.
- With a maximum limit of 15 files, these set of created files would be passed in as arguments to `7zz`.

We created a new `MainAr.cpp` with the `main` function being replaced by our harness which would mutate the input and pass it into `argv` of `main_7zz`, the original `main` function of `p7zip`. This way, we can fuzz the archive command with multiple files and potentially discover more bugs.

```
1 AFL_SKIP_CPUFREQ=1 AFL_I_DONT_CARE_ABOUT_MISSING_CRASHES=1 $(AFL_FUZZ)
  -M main-afl-$(HOSTNAME) -t 30000 -i in -o out -- $(BIN_AFL_HARNESS)
  @@
```

Our `afl-fuzz` command just passes in a cli argument to the harness.

Results

TODO

OSS-Fuzz and State of Fuzzing p7zip

When we were researching good ways to fuzz `p7zip`, we found a [pull request](#) on the repository that asks the maintainer if they would like `p7zip` to be part of a wider fuzzing initiative by Google called

[OSSFuzz](#) which fuzzes popular open-source projects for free. Although this pull request was never merged, it hinted to us that it was possible that some organizations are fuzzing this project regularly.

Interestingly, we also found that the repository itself has a fuzzer set up for some of its codecs (compression algorithms).

For example, we found a dependency [zstd](#), which is a real-time compression library build by Meta. This has a fuzzing subdirectory set up for fuzzing this library.

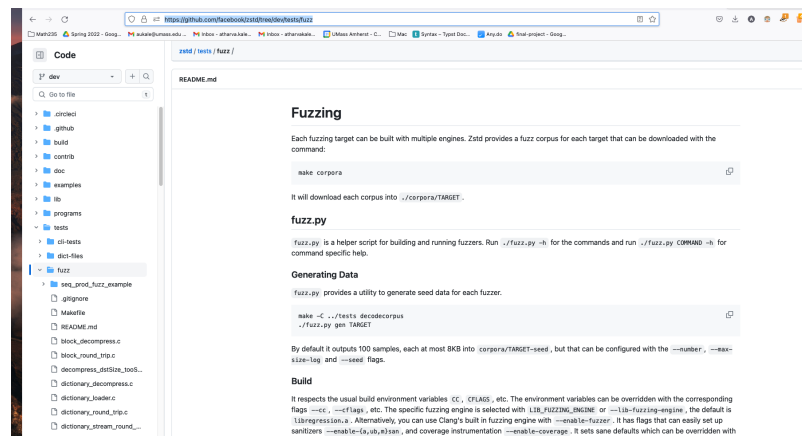


Figure 7: zstd Fuzzing Setup

We found that they are using [libFuzzer](#), a coverage guided fuzzing library by LLVM. They also use AFL in tandem with their setup.

We found really interesting design choices in their fuzzing setup. One such choice was to incorporate unit-fuzzing, which is fuzzing small features of the target.

This was evident in a simple LLVM harness they wrote:

```
zstd / tests / fuzz / simple_compress.c

Code Blame 60 lines (51 loc) · 1.87 KB

21 #include "zstd.h"
22 #include "zstd_errors.h"
23 #include "zstd_helpers.h"
24 #include "fuzz_data_producer.h"
25 #include "fuzz_third_party_seq_prod.h"
26
27 static ZSTD_CCtx *cctx = NULL;
28
29 int LLVMFuzzerTestOneInput(const uint8_t *src, size_t size)
30 {
31     FUZZ_SEQ_PROD_SETUP();
32
33     /* Give a random portion of src data to the producer, to use for
34     parameter generation. The rest will be used for (de)compression */
35     FUZZ_dataProducer_t *producer = FUZZ_dataProducer_create(src, size);
36     size = FUZZ_dataProducer_reserveDataPrefix(producer);
37
38     size_t const maxSize = ZSTD_compressBound(size);
39     size_t const bufSize = FUZZ_dataProducer_uint32Range(producer, 0, maxSize);
40
41     int const cLevel = FUZZ_dataProducer_int32Range(producer, kMinCLevel, kMaxCLevel);
42
43     if (!cctx) {
44         cctx = ZSTD_createCCtx();
45         FUZZ_ASSERT(cctx);
46     }
47
48     void *rBuf = FUZZ_malloc(bufSize);
49     size_t const ret = ZSTD_compressCCtx(cctx, rBuf, bufSize, src, size, cLevel);
50     if (ZSTD_isError(ret)) {
51         FUZZ_ASSERT(ZSTD_getErrorCode(ret) == ZSTD_error_dstSize_tooSmall);
52     }
53     free(rBuf);
54     FUZZ_dataProducer_free(producer);
55 #ifndef STATEFUL_FUZZING
56     ZSTD_freeCCtx(cctx); cctx = NULL;
57 #endif
58     FUZZ_SEQ_PROD_TEARDOWN();
59     return 0;
60 }
```

Figure 8: Simple Compress Harness

In this harness, they first produce an input using the libFuzzer's input generator. They then call into the function they are fuzzing-`ZSTD_compressCCtx`.

```
262
263  /*! ZSTD_compressCctx() :
264    * Same as ZSTD_compress(), using an explicit ZSTD_Cctx.
265    * Important : in order to behave similarly to `ZSTD_compress()`,
266    * this function compresses at requested compression level,
267    * __ignoring any other parameter__ .
268    * If any advanced parameter was set using the advanced API,
269    * they will all be reset. Only `compressionLevel` remains.
270    */
271  ZSTDLIB_API size_t ZSTD_compressCctx(ZSTD_Cctx* cctx,
272                                       void* dst, size_t dstCapacity,
273                                       const void* src, size_t srcSize,
274                                       int compressionLevel);
275
```

Figure 9: ZSTD_compressCctx()

This seems to be a major function in their library that handles compression of certain input frames.

These more advanced approaches seem more well suited to a project of this size, as fuzzing with a small harness may provide a higher exec speed.

Static Analysis

Challenges Faced

Working with a large C/C++ codebase

It was our first exposure to working with a large C/C++ codebase. Although neatly organized at first glance, the project quickly turned into a codebase with a bunch of build scripts. Our first challenge was to figure out how to get a debug and release build going. Documentation on the dependencies was sparse, so this involved a compile-and-fail cycle to find all the dependencies for our systems. However, this experience provided us with great insight on how real world C++ projects are build, and gave us some direction on how to design such a codebase for a project in the future.

Real world projects

Next Steps